AD CLICK PREDICTION USING MACHINE LEARNING

A PROJECT REPORT

Submitted by

Dimpy Garg(22BAI71408)
Kanika (22BAI71383)
Bhuvana(22BAI71358)
Udaya Sri (22BAI71389)

in partial fulfilment for the award of the degree of

Bachelor of Engineering

IN

Artificial Intelligence and Machine Learning

Chandigarh University

APRIL 2026

Chandigarh University

BONAFIDE CERTIFICATE

Certified that this project report " **Ad - click prediction using machine learning**" is the bonafide work of "Dimpy Garg, Kanika , Bhuvana , Udaya Sri " who carried out the project work under my/our supervision.

……. ……..

SIGNATURE SIGNATURE

Mr. MA Kumar Dr. Suresh Kaswan

Supervisor HOD

submitted for the project viva-voce examination held on

……. ……..

INTERNAL EXAMINER EXTERNAL EXAMINER

Acknowledgement

It grants us the opportunity to finish this midterm assignment. This is the only page where I can express our feelings of gratitude and emotion. It gives me great pleasure to extend my sincere gratitude to **Dr. Manchu Abhilash Kumar ,** my supervisor and mentor, for all of her helpful advice, direction, and unwavering support during the completion of this project, **"AD-Click Prediction using Machine Learning."**Even though we are the ones working on this project, we could not have done it without the assistance of many others, whose collaboration has enabled us to complete it successfully. I am really very thankful to Chandigarh University for providing us with such a great opportunity to make such a wonderful project that can solve real-life problems and provides extremely valuable hands-on experience along with crucial soft skills such as working in a team, communication skills, and much more. I also offer my most sincere thanks to every team member of our group who was working vigorously on this project and staff members of Apex Institute of Technology, Chandigarh University, for the cooperation provided by them in every possible way. We thank all the faculty members and other supporting staff for the help they provided to us for the completion of our project.

Dimpy Garg (22BAI71408)

Kanika (22BAI71383)

Bhuvana (22BAI71358)

Udaya Sri (22BAI71389)

Table of content

Abstract.....7

CHAPTER 1 . INTRODUCTION

1. Identification of Client / Need / Relevant Contemporary issue.....8
2. Identification of Problem.....9
3. Identification of tasks.....10
4. Timeline… 11
5. Organisation of Report 12

CHAPTER 2. LITERATURE REVIEW

1. Timeline.....13
2. Existing solutions.....14
3. Bibliometric Analysis 16
4. Review Summary… 17
5. Problem Definition 18

CHAPTER 3.DESIGN FLOW AND PROCESS

1. Evaluation and Selection of Process 19
2. Design Constraints.....20
3. Analysis of Feature Finalization.....21
4. Design Flow and Implementation… 22
5. Design Selection.....23
6. Methodology 24

CHAPTER 4. RESULT ANALYSIS AND VALIDATION

1. Implementations Of Design Using Modern Engineering Tools In Analysis ……35
2. Result and Testing.....40

3. Report Preparation: 42

4. Project management.....44

CHAPTER 5. CONCLUSION AND FUTURE WORK

1. Conclusion.....46

2. Future Work.....47

List of Figures

Figure 1 11
Figure 2 14
Figure 3 38
Figure 4 38
Figure 5 38
Figure 6 38
Figure 7 39
Figure 8 39

ABSTRACT

This report delves into an investigation of the efficacy of logistic regression as a predictive modelling technique for ad click prediction. The study seeks to explore the performance of logistic regression and assess its predictive capabilities through a range of evaluation metrics. The primary goal is to assess the accuracy, precision, recall, F1 score, and area under the receiver operating characteristic curve (AUC-ROC) of logistic regression models in the context of predicting ad clicks. The paper commences with an extensive review of existing literature on ad click prediction and logistic regression, emphasising the significance of precise click-through rate (CTR) estimation for advertisers and the potential advantages of utilising logistic regression in this domain. The theoretical foundation of logistic regression is also addressed, providing a comprehensive understanding of its underlying principles and assumptions. Subsequently, the methodology section outlines the dataset employed for the study, incorporating historical ad impressions and click data. The steps involved in data pre-processing, encompassing feature selection and transformation, are elucidated. Logistic regression models are then trained on the pre-processed data, and their performance is assessed using diverse evaluation metrics. The results section divulges the study's findings, presenting a detailed analysis of accuracy, precision, recall, F1 score, and AUC-ROC derived from the logistic regression models. Model performance is benchmarked against alternative algorithms commonly employed in ad-click prediction. The results underscore the strengths and limitations of logistic regression in predicting ad clicks. Furthermore, the discussion section provides insights into the implications of these findings, exploring the interpretability of logistic regression models and their potential to offer actionable insights to advertisers. The section also addresses the limitations and potential challenges associated with using logistic regression in ad click prediction. Finally, the conclusion section succinctly summarises the key research findings and provides recommendations for future studies, emphasising the reliability and interpretability of logistic regression for ad click prediction,

CHAPTER 1
INTRODUCTION

1. . Identification of Client /Need / Relevant Contemporary issue

Logistic Regression is commonly used to estimate the probability that an instance belongs to a particular class. If the estimated probability that an instance is greater than 50%, then the model predicts that the instance belongs to class 1, or else it predicts that it does not. This makes it a binary classifier.

Data visualisation is a crucial aspect of machine learning that enables analysts to understand and make sense of data patterns, relationships, and trends. Data visualisation helps machine learning analysts better understand and analyse complex data sets by presenting them in an easily understandable With the increasing availability of big data, it has become more important than ever to use data visualisation techniques to explore and understand the data.

In this project, we are analysing the performance and efficiency of logistic regression by taking an example of ad click prediction.

There are various advertisements in our daily lives; some of them are useful and attractive to us, while others are useless to us, and we will never click on them. As an advertiser, our aim is to maximise the effectiveness of your advertising campaigns by understanding the factors that influence user engagement with our ads. In today's digital age, online advertising has become a cornerstone of marketing strategies across industries. However, with the vast array of advertisements bombarding users daily, the challenge for advertisers lies in ensuring their ads capture the attention of the right audience at the right time. Ad click prediction addresses this challenge by leveraging data-driven techniques to forecast the likelihood of a user clicking on an advertisement. With the proliferation of online platforms and the vast amount of content vying for users' attention, advertisers face the challenge of ensuring that their ads are not only seen but also interacted with by the right audience. Ad click prediction addresses this challenge by analysing various factors that influence user behaviour, such as demographics, browsing history, and contextual information.

From the perspective of the user, it is disgusting to receive many useless advertisements when surfing the Internet. From the perspective of the advertiser, it is depressing that the given advertisement is not attractive and receives an extremely low click rate.

In this project, we are going to work on an advertising dataset, indicating whether or not a particular internet user has clicked on an advertisement.

Ultimately, the goal of ad click prediction is to empower advertisers with actionable insights that enable them to target their ads more precisely, increase click-through rates, and maximise the return on investment from their advertising efforts.

2. Identification Of Problem :

In the dynamic landscape of advertising, where capturing consumer attention is crucial, predicting user engagement with advertisements becomes imperative for maximizing campaign effectiveness and ensuring a high return on investment. Various predictive models and techniques have emerged to tackle this challenge, each offering its own set of advantages and capabilities. However, understanding how these models perform in comparison to each other is essential for advertisers to make informed decisions about their advertising strategies. In this study, our main objective is to evaluate the performance of logistic regression, a widely-utilized statistical technique, in predicting ad clicks. We aim to assess logistic regression's ability to forecast whether a user will click on an advertisement based on key features such as demographics, browsing behavior, and contextual information. Through rigorous experimentation and evaluation using real-world advertising datasets, we intend to quantify logistic regression's predictive performance, compare it against other machine learning algorithms, and analyze its strengths and limitations in ad click prediction tasks. By shedding light on the effectiveness of logistic regression and its potential to provide actionable insights for advertisers, we seek to empower businesses to optimize ad targeting strategies and enhance click-through rates. Ultimately, this study aims to contribute to the advancement of ad click prediction methodologies by offering a comprehensive assessment of logistic regression's performance, thereby enabling advertisers to make informed decisions and drive success in the competitive realm of digital advertising.

3. IDENTIFICATION OF TASKS

1. Dataset Collection:

- The goal of dataset collection is to obtain a comprehensive and representative set of data that can be used for analysis and predictive modelling.
- For this project, the dataset comprises five variables: open, close, low, high,

and volume. These variables represent different bid prices for a stock at different times and the volume of shares traded during those periods.

- The dataset will be collected from reliable sources such as financial databases or APIs, ensuring accuracy and relevance to the problem domain.
- Once collected, the dataset will be divided into training and testing sets, with a portion reserved for model evaluation.

2. Processing the Dataset:

- The collected data will undergo preprocessing and formatting to prepare it for analysis and modelling.
- Various steps will be taken to clean the data, handle missing values, and transform variables as needed.
- Feature engineering techniques may be applied to extract additional information or create new features that could improve model performance.
- The processed dataset will be structured in a way that facilitates the desired output, making it suitable for predictive modelling.

3. Comparing different ML Algorithms/Literature Review:

- The next objective is to apply different machine learning algorithms to the processed dataset.
- These algorithms may include logistic regression, decision trees, random forests, support vector machines, and neural networks, among others.
- The performance of each algorithm will be evaluated based on metrics such as accuracy and recall.
- A literature review will also be conducted to explore existing research and approaches in the field of click prediction, providing insights and guidance for the project.

4. Problem Formulation:

- Once the data is formatted, various graphical insights will be derived from the dataset.
- This could include visualisations of stock price trends over time, comparisons between different companies, and correlations between variables.
- Confusion matrices and similar insights will be created to facilitate the visualisation and interpretation of model performance.
- These visualisations will help in understanding the underlying patterns and relationships in the data, guiding further analysis and model development.

5. Selecting the suitable ML model:

- After comparing the performance of different machine learning models,

the most suitable model will be selected based on its accuracy and other relevant metrics.

- The selected model will be one that demonstrates the best predictive performance for determining whether a viewer will click on an advertisement.
- Consideration will also be given to the model's interpretability, scalability, and ease of implementation in real-world scenarios.

6. Testing and training the dataset:

- The selected model will be trained on the training dataset using appropriate techniques, such as cross-validation, to optimise its parameters.
- Various evaluation metrics will be used to assess the model's performance on the testing dataset, ensuring its generalizability and robustness.
- Hyperparameter tuning may be performed to further improve the model's accuracy and performance.

7. Predicting the rate of ad clicks:

- Once the model is trained and validated, it will be used to predict the likelihood of a viewer clicking on an advertisement.
- These predictions will provide valuable insights for companies seeking to customise their ads and optimise their marketing strategies.
- By leveraging the predictive power of the model, companies can enhance their sales and productivity by targeting ads more effectively at their audience.

1.5 Timeline

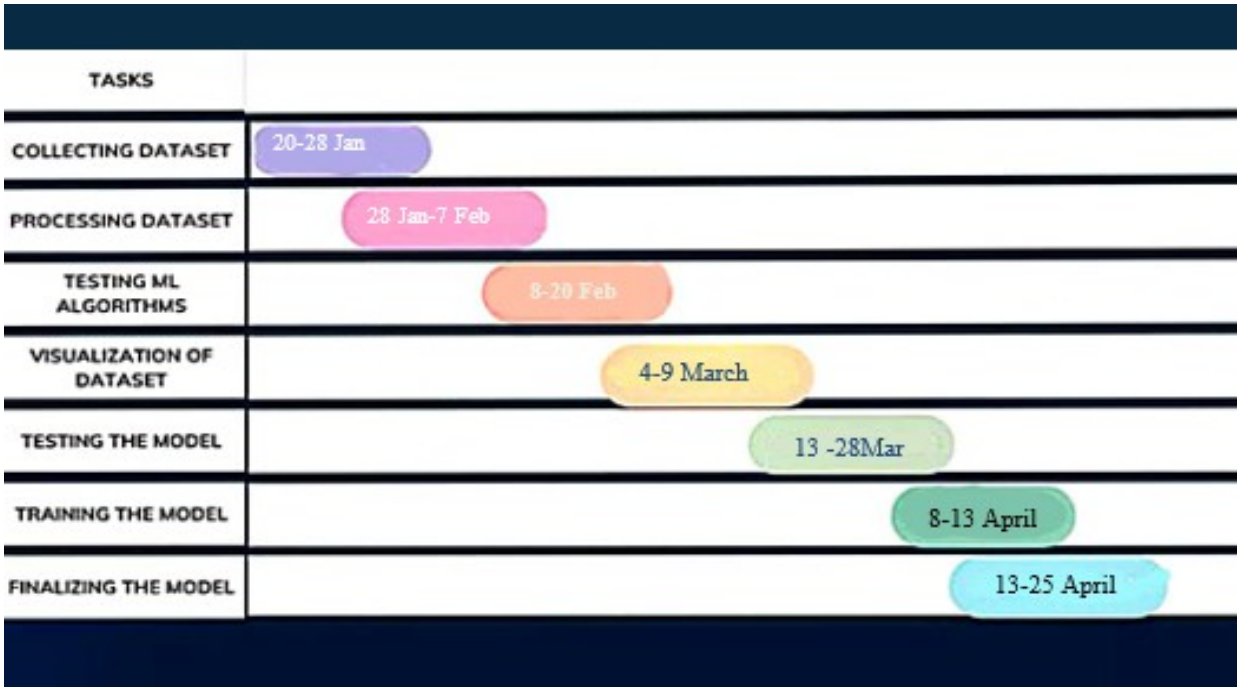

Fig. 1

1.5 Organization of the Report

Chapter 1 Problem Identification: This chapter introduces the project and describes the problem statement discussed earlier in the report

Chapter 2 Literature Review: This chapter prevents review for various research papers which help us to understand the problem in a better way. It also defines what has been done to already solve the problem and what can be further done.

Chapter 3 Design Flow/ Process: This chapter presents the need and significance of the proposed work based on literature review. Proposed objectives and methodology are explained. This presents the relevance of the problem. It also represents a logical and schematic plan to resolve the research problem.

Chapter 4 Result Analysis and Validation: This chapter explains various performance parameters used in implementation. Experimental results are shown

in this chapter. It explains the meaning of the results and why they matter.

Chapter 5 Conclusion and future scope: This chapter concludes the results and explains the best method to perform this research to get the best results and define the future scope of study that explains the extent to which the research area will be explored in the work.

CHAPTER 2

LITERATURE REVIEW / BACKGROUND STUDY

1. Timeline of the reported problem

"On October 27, 1994, the world of advertising was forever transformed by a small graphic bearing the presumptive words, "Have you ever clicked your mouse right here? You will," in a kitschy rainbow font. The age of banner ads had officially begun."

"In 1996, banner ads plastered the internet, but advertisers still didn't have a good process to determine if these ads were actually driving tangible results for their businesses. Marketers needed a way to more efficiently manage their display ad campaigns across multiple websites and report on how users were interacting with their ads."

"It would be an understatement to say that pop-up ads suffer from a poor image problem. They've been called the internet's original sin and the most hated advertising technique, and one of the original developers has even apologised for creating the underlying code that unleashed them upon unsuspecting web surfers. Even so, these much-maligned ads hold an undeniable place in the history of online advertising."

"By the time 1999–2002, the web was expanding rapidly, and users needed a better way to navigate the terrain. With search engines steadily gaining popularity, advertisers looking to create ads that were more targeted and less loathsome turned to sponsored search as the next digital advertising frontier."

"Around 2010–present, a new group of media companies began to emerge. Websites like BuzzFeed and Mashable presented advertisers with new opportunities to connect with their audiences through sponsored content and native advertising."

"While the goal of the content is to promote the business, the format looks and feels less like an ad and more like a regular piece of content on the host's website."

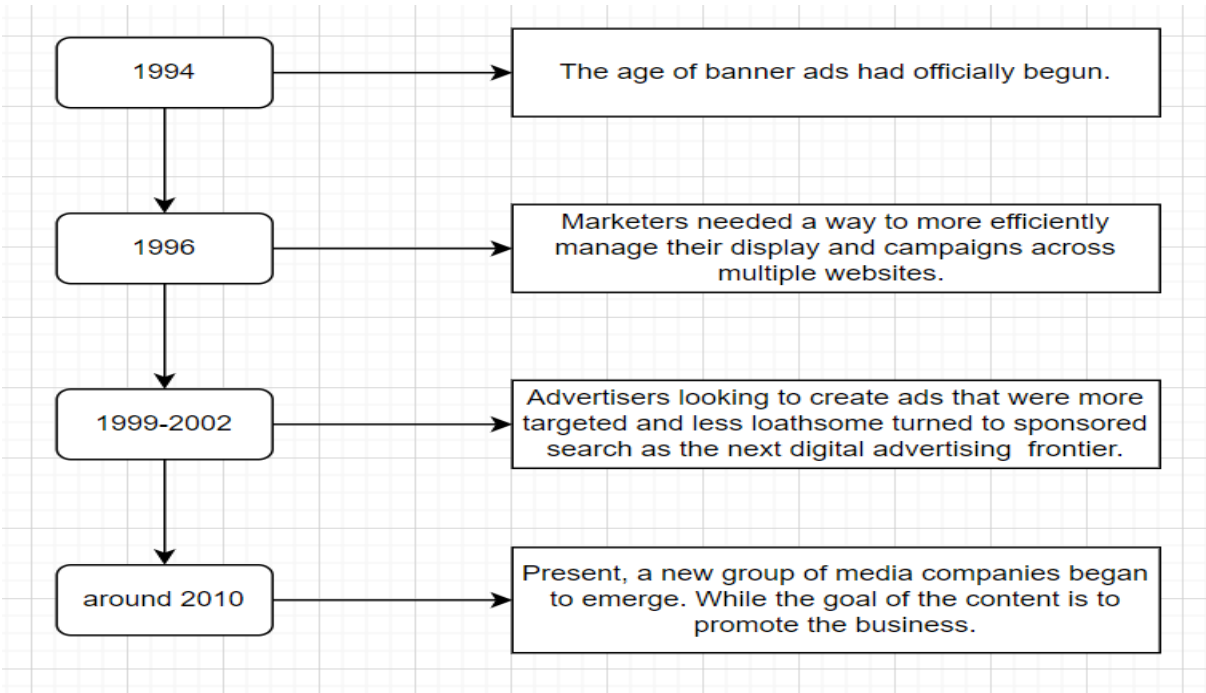

Fig.2

2. Existing Solutions and Review Summary

A number of research works to find solution to the ad-click prediction have already been done using various machine learning techniques like Logistic regression and Naive Baye's

Classifiers, support vector machines, and decision trees. Logistic regression has a central role in the early analysis of this domain.

From a purely statistical standpoint, it depends—on the campaign one is running, the keywords targeting, and the channel and purpose of the ad. So while one wants to have a "high" click-through rate, there's really no magic number. According to the survey for attorneys and legal services, the average click-through rate is 3.85%, so a good click-through rate for this industry might be something like 5-6%. For arts and entertainment, the average click-through rate is 10.67%, so a good CTR for businesses in this industry would be something like 11–12%.

Several studies have explored the effectiveness of different machine learning algorithms in understanding user behaviour and predicting ad clicks.

Logistic regression: Research has shown its potential in uncovering user behaviour and targeting advertisements. For instance, a Twitter advertising prototype successfully employed logistic regression.

Naive Bayes: This method has also played a significant role in ad-click prediction, as evidenced by various studies.

Decision trees: Combined with logistic regression, decision trees have proven effective in serving targeted ads on social media platforms like Facebook.

With the continuous rise of online activity, analysts are actively seeking new methods to display relevant advertisements that align with user interests. Research has increasingly focused on leveraging machine learning for ad targeting, particularly on large scales. The recent boom in social networking has provided real-world examples of successful ad-serving mechanisms like those employed by Twitter and Facebook.

A multitude of research endeavours have been undertaken to address the challenge of ad-click prediction, employing a diverse array of machine learning techniques such as Logistic Regression, Naive Bayes Classifier, Support Vector Machines, and Decision Trees. Among these, Logistic Regression has emerged as a cornerstone in the early exploration of this domain. However, determining an optimal click-through rate is not a straightforward endeavour, as it varies depending on factors such as the nature of the campaign, targeted keywords, and the channel and purpose of the ad. While advertisers aim for a high click-through rate, there is no universal "magic number." For instance, in the legal services industry, the average click-through rate hovers around 3.85%, suggesting that a good click-through rate for this sector might range from 5-6%. Conversely, the arts and entertainment industry boasts a higher average click-through rate of 10.67%, indicating that a satisfactory CTR for businesses in this domain could be around 11–12%.

Numerous studies have delved into the efficacy of various machine learning algorithms in deciphering user behaviour and predicting ad clicks. Logistic regression, for example, has been extensively studied for its potential to uncover intricate patterns in user behavior and facilitate targeted advertisement delivery. Notably, a Twitter advertising prototype successfully leveraged logistic regression to enhance ad targeting and engagement. Similarly, the Naive Bayes method has played a significant role in ad-click prediction, as evidenced by its utilization in numerous studies exploring user behavior and ad engagement. Additionally, the combination of decision trees with logistic regression has proven effective in serving targeted ads on prominent social media platforms like Facebook.

As online activity continues to surge, analysts are increasingly exploring innovative approaches to deliver relevant advertisements that resonate with user interests. Research endeavors have progressively shifted towards harnessing the power of machine learning for ad targeting, particularly at scale. The proliferation of social

networking platforms has provided tangible examples of successful ad-serving mechanisms, exemplified by the strategies employed by industry giants such as Twitter and Facebook. By leveraging machine learning algorithms and vast datasets, advertisers can enhance the relevance and effectiveness of their advertisements, ultimately improving user engagement and driving business growth in the competitive landscape of online advertising.

Review Summary:

Aspect	Summary
Goal	Estimate the probability of a user clicking on an ad.
Benefits	- Optimize campaign performance by allocating budget to high-click ads. - Improve targeting by tailoring ads to receptive user segments. - Maximize ROI by increasing conversions from clicks.
Logistic Regression	Pros: Simple, interpretable, efficient. Cons: May not capture complex non-linear relationships.
Data Preparation	- Feature Engineering (user, ad, context features). - Data Cleaning (outliers, missing values). - Data Splitting (training, validation, testing).
Model Training	Train the model on the training data. Tune hyperparameters for optimal performance.
Evaluation	Accuracy, precision, recall, F1-score, AUC-ROC curve.
Data Visualization	Pair plots, heatmaps, distribution plots, interactive dashboards. Benefits: Explore relationships, identify patterns, understand model behavior. Drawbacks: Limited scope, human bias.
Solutions	- Ensemble Methods: Combine logistic regression with other models (e.g., decision trees) to handle non-linearities. - Feature Importance: Analyze coefficients or use SHAP values to identify impactful features. - Real-Time Scoring: Optimize models for real-time click prediction.
Integration	- Data Warehouses/Lakes (historical data access). - Campaign Management Platforms (real-time ad serving). - Data Visualization Tools (deeper click prediction insights).

Additional Considerations	- Data Quality (clean, structured data). - Model Selection (compare logistic regression with other algorithms). - Continuous Learning (regular updates with new data).

3. Biblometric analysis

Bibliometric analysis involves the quantitative analysis of scholarly literature, including publication patterns, citation networks, and the impact of research outputs. In the context of our study on ad click prediction, bibliometric analysis provides

insights into the existing body of literature on similar topics, identifies key researchers and institutions, and highlights emerging trends and research gaps.

Key Features:

In our bibliometric analysis, we identify key features that are commonly examined in studies related to ad click prediction. These features provide valuable insights into user behavior, preferences, and characteristics that influence ad engagement. Let's
explore some of these key features along with data examples:

Age: (Categorical)

Users in different age groups may exhibit varying levels of receptiveness to different types of ads. For example, users in the 25-34 age group might be more receptive to
ads for video games compared to users over 55 who may prefer ads for retirement planning or travel. By categorizing users based on age demographics, advertisers can tailor their ad campaigns to target specific age groups more effectively.

Time Spent on Platform: (Numerical)

The amount of time users spend on a social media platform can also impact their likelihood of clicking on ads. For instance, users who spend more than 30 minutes daily on a social media platform are likely to have a higher level of engagement and may be more receptive to ads compared to those who log in briefly. Advertisers can use this information to adjust their ad targeting strategies and allocate resources more efficiently.

Internet Usage: (Categorical)

The devices users use to access the internet can influence their ad-clicking behavior. For example, users who primarily access the internet from mobile devices might be more susceptible to clicking on mobile app install ads compared to desktop users.
Understanding users' preferred devices can help advertisers tailor their ad formats and placements to maximize engagement and conversion rates.

Salary: (Numerical)

Income level can also play a role in ad engagement, as individuals with higher incomes may have different purchasing behaviors and preferences. For example, individuals with higher incomes may be more receptive to ads for luxury goods
or investment products. Advertisers can use income data to target high-income
demographics more effectively and tailor their ad content to appeal to
their preferences and lifestyles.Data Preprocessing with Specific Examples:

In our bibliometric analysis, we employ various data preprocessing techniques to clean and prepare the data for analysis. Let's examine some specific examples of data preprocessing techniques used in our study:

Label Encoding:

One common preprocessing technique is label encoding, which involves converting categorical variables into numerical format. For example, a binary target variable such as "click" or "no click" can be encoded, where "click" is assigned a value of 1 and "no click" is assigned a value of 0. This allows us to incorporate categorical variables into our predictive models more effectively.

Missing Value Imputation:

Another preprocessing technique is missing value imputation, which involves filling in missing data points in the dataset. For example, if there are missing data points in features like "Time Spent," we can fill them in with the median value for that specific user category. This ensures that our analysis is based on complete and reliable data, improving the accuracy and robustness of our findings.

By employing these data preprocessing techniques, we ensure that our bibliometric analysis is based on clean, reliable, and informative data, allowing us to derive meaningful insights into the trends, patterns, and relationships in scholarly literature related to ad click prediction.

Model Training and Testing with Performance Metrics:

- **Split Ratio:**

  The data can be split into an 80/20 training/testing ratio, but a more balanced split (e.g., 70/30) might be used for smaller datasets.

- **Evaluation Metrics:**

  Beyond accuracy, metrics like AUC-ROC (Area Under the Receiver Operating Characteristic Curve) can assess how well the model distinguishes between clickers and non-clickers.

1. **Data Preparation:**
   - **Feature Engineering**:
     - ❖ Create informative features (e.g., time spent on similar categories for predicting ad click interest).
     - ❖ Handle categorical features (e.g., one-hot encoding).
     - ❖ Address missing values (e.g., imputation techniques).
   - **Data Cleaning:**
     - ❖ Identify and remove outliers or inconsistencies.
   - **Data Splitting:**
     - ❖ Divide data into training, validation (for hyperparameter tuning), and testing sets (to assess model performance.
2. **Model Training:**
   - Train the logistic regression model using the training data.
   - Tune hyperparameters (e.g., regularisation) to optimise performance.
3. **Evaluation:**
   - Evaluate the model's performance on the testing data using metrics like accuracy, precision, recall, F1-score, and AUC-ROC curve.

Data Visualization for Enhanced Insights

Visualisation plays a crucial role in understanding click behaviour and improving model performance:

- ❖ **Pair Plots:**

  Visualise relationships between pairs of features (e.g., user age vs. click-through rate) to identify potential correlations.

- **Heatmaps:**

  Explore relationships between multiple features to uncover patterns in click behaviour.

- **Distribution Plots:**

  Analyse distributions of individual features (e.g., user demographics, ad categories) to identify potential biases.

- **Interactive Dashboards:**

  Create interactive tools to explore and filter data dynamically, aiding in identifying feature interactions and model shortcomings.

Addressing Drawbacks with Solutions:

- **Model Limitations:**

  Consider ensemble methods that combine predictions from multiple logistic regression models or explore more advanced algorithms like gradient boosting for potentially better performance on complex relationships.

- **Data Quality Dependency:**

  Implement data validation checks to identify and address anomalies or inconsistencies. Regularly monitor model performance and retrain with fresh data to account for evolving user behaviour.

- **Ensemble Methods:**

  Combine multiple models (e.g., decision trees, random forests)
  to handle non-linearities.

- **Feature Engineering:**

  Create new features that capture non-linear interactions.

- **Interactive Visualization:**

  Use interactive dashboards to explore data dynamically and uncover hidden patterns.

- **Model-Agnostic Visualizations:**

  Explore SHAP values (SHapley Additive exPlanations) for feature importance insights, regardless of the model type.

- **Sensitive to Outliers:**

  Logistic regression is sensitive to outliers in the data, which can skew parameter estimates and affect the performance of the model.

- **Interpretability Challenges:**

  Utilise feature importance scores to understand which features have the strongest influence on the model's predictions.

- **Potential Misuse:**

  Develop ethical guidelines for data collection, model training, and ad targeting practices. Implement mechanisms to detect and prevent click fraud.

- **Data Imbalance:**

  Clicks are often rarer than non-clicks. Use techniques like SMOTE (Synthetic Minority Oversampling Technique) to handle imbalance.

- **Model Explainability:**

  Employ LIME (Local Interpretable Model-Agnostic Explanations) to gain human-understandable explanations from complex models.

- **Real-Time Updates:**

  Update models continuously with new data for consistent accuracy.

4. Problem Definition

In the realm of advertising, the ability to predict user engagement with advertisements is paramount for maximising campaign effectiveness. Various predictive models exist for this purpose, each offering unique strengths and capabilities. In this study, we aim to evaluate the performance of logistic regression —a widely-used statistical technique—in predicting ad clicks.Specifically, our goal is to assess how well logistic regression performs in forecasting whether a user will click on an advertisement based on relevant features such as demographics, browsing behaviour, and contextual information. By comparing the predictive accuracy, interpretability, and computational efficiency of logistic regression against other models, we seek to determine its suitability for ad click prediction tasks.

Through rigorous experimentation and evaluation on real-world advertising datasets,

we will quantify the predictive performance of logistic regression and identify its strengths and limitations. This analysis will provide advertisers with valuable insights into the efficacy of logistic regression as a predictive tool for optimising ad targeting and maximising click-through rates.

Overall, this study aims to contribute to the advancement of ad click prediction methodologies by offering a comprehensive assessment of logistic regression's performance in this domain. By elucidating the capabilities of logistic regression, advertisers can make informed decisions about the choice of predictive models and enhance the effectiveness of their advertising strategies.

5. Goals / Objectives

- The goal is to create and train machine learning models, such as logistic regression, to predict ad clicks.
- **Optimize Campaign Performance:** Allocate budgets more effectively by focusing on ads with a higher likelihood of clicks.
- **Improve Targeting:** Tailor ads to specific user segments who are more receptive to the message.
- **Maximize ROI (Return on Investment):** Increase the number of conversions (e.g., purchases, sign-ups) generated from ad clicks.
- **High Accuracy:** The model should accurately distinguish between instances where a user will click (positive class) and those where they won't (negative class). Common metrics for accuracy include:
    - **Accuracy:** Proportion of correctly classified clicks and non-clicks.
    - **AUC-ROC:** Area Under the Receiver Operating Characteristic Curve, measures performance across all classification thresholds.
- **Interpretability:** The model's internal workings should be understandable to some degree. Logistic regression offers an advantage here as coefficients associated with features provide insights into their influence on click probability.
- **Feature Importance:** Identify which user, ad, and contextual features have the most significant impact on click behavior. This information can be used for further feature engineering and campaign optimization.
- **Generalizability:** The model should perform well on unseen data (new users, ad formats, etc.). This can be achieved through techniques like cross-validation and regularization to avoid overfitting.

- **Scalability:** The model should be able to handle large datasets efficiently and adapt to changing advertising landscapes. Efficient algorithms and cloud-based computing can be leveraged for scalability.
- **Real-Time Performance:** For certain applications, it might be desirable to predict clicks in real-time for ad serving optimization. This can be achieved through model optimization and efficient deployment techniques.
- **Data Quality:** The success of your model heavily relies on the quality of your data. Ensure you have a clean, well-structured dataset with minimal missing values and outliers.
- **Model Selection:** While logistic regression can be effective, it's important to consider other machine learning algorithms and compare their performance to find the best fit for your specific data and objectives.
- **Continuous Learning:** Regularly update your model with fresh data to maintain its accuracy and adapt to evolving user behavior and advertising trends.
- The project will have several milestones, including collecting and processing data, feature engineering, model development, and evaluation. The feature engineering process will include identifying the most relevant features for stock market prediction, performing feature selection and extraction, and scaling and normalizing the data.
- Develop and implement a machine learning model using SVM and decision tree algorithms to predict the stock prices of a specific company with at least 80% accuracy, validated through backtesting.
- Analyse and compare the performance of the logistic regression model in predicting ad clicks Analyze, identifying the strengths and weaknesses of each approach.
- Analyse the influence of various features, including user demographics, device type, time of day, ad content, and historical click-through rates, on the likelihood of ad clicks using logistic regression.
- Evaluate the potential profitability of an ad placement strategy based on predicted click probabilities. Simulate the performance of the strategy using historical ad impressions and click data over a specific period to estimate potential revenue gains.
- Conduct sensitivity analyses to assess the robustness of the logistic regression model and the ad placement strategy under different market conditions. Evaluate performance metrics such as accuracy, precision, recall, and F1-score across varying levels of ad inventory, competition, and user engagement.

CHAPTER 3.

DESIGN FLOW AND PROCESS

1. Evaluation and Selection of Features

In this section, we embark on the critical task of identifying, evaluating, and selecting features that will be used to train our logistic regression model for click prediction.

1. **Feature Identification**:

   We initiate the process by brainstorming and researching potential features that might influence the likelihood of a user clicking on an ad. These features could range from user-specific attributes like demographics and browsing history to ad-specific characteristics such as size and placement.

2. **Feature Exploration**:

   Once potential features are identified, we meticulously explore their distributions and relationships with the target variable (click or no-click). We employ various visualisation techniques, such as histograms, scatter plots, and heatmaps, to gain insights into the significance of each feature and its potential predictive power.

3. **Feature Selection Methods:**

   To determine the most relevant features for our logistic regression model, we leverage a combination of feature selection methods. This includes:

   - Correlation analysis: identifying features that exhibit high correlation with the target variable. Features with strong correlations are likely to have a significant impact on click prediction.
   - Feature importance: utilizing techniques such as information gain or feature importance scores from ensemble models like Random Forest or Gradient Boosting Machines. These methods help rank features based on their predictive power.
   - Domain knowledge: drawing upon domain expertise to handpick features that are known to influence user click behaviour based on prior research or industry insights.

4. **Cross-validation:**

   To ensure the robustness of our feature selection process, we employ cross-validation techniques such as k-fold cross-validation. This involves splitting the data into multiple folds, training the model on different subsets, and evaluating its performance across each fold to guard against overfitting.

2. Design Constraints

In this section, we identify and address various constraints that may influence the design and implementation of our click prediction system.

1. **Data Availability**:

   We acknowledge potential limitations in the availability, quality, and diversity of data required for training our logistic regression model. This may include challenges such as missing values, imbalanced classes, or insufficient sample size, which could impact the model's predictive performance.

2. **Computational Resources**:

   We consider the computational resources needed for training and deploying the logistic regression model, especially for large-scale datasets. Factors such as memory requirements, processing speed, and scalability are taken into account to ensure efficient model training and inference.

3. **Regulatory Compliance:**

   Compliance with data privacy regulations such as GDPR (General Data Protection Regulation) or CCPA (California Consumer Privacy Act) is paramount. We adhere to legal and ethical guidelines governing data collection, storage, and processing to safeguard user privacy and maintain regulatory compliance.

3. Analysis of Feature Finalization

With a thorough understanding of feature relevance and constraints, we finalize the set of features to be used in our logistic regression model.

1. **Feature Engineering:** We undertake necessary preprocessing steps to refine and engineer features for optimal model performance. This includes:
   - **Handling missing data**:

Employing techniques such as mean imputation, median imputation, or predictive imputation to address missing values.

- **Encoding categorical variables**:
  Converting categorical variables into numerical representations through one-hot encoding, label encoding, or target encoding.
- **Feature scaling:**
  Normalizing or standardising numerical features to ensure they are on a similar scale, preventing certain features from dominating others during model training.

2. **Feature Selection:**

   Based on our analysis and domain expertise, we curate a subset of features that demonstrate the highest predictive power for click prediction. We prioritise features that exhibit strong correlations, high importance scores, or strong theoretical justification for inclusion in the model.

3. **Domain Relevance**:

   Features should be selected based on their relevance to the domain and the specific task of ad-click prediction. Features that directly relate to user behavior, content characteristics, and contextual factors influencing ad engagement are likely to have higher predictive power.

4. **Dimensionality Reduction:**

   In cases where the feature space is large, employing techniques such as principal component analysis (PCA) or feature importance ranking can help reduce dimensionality while preserving the most informative features. This can improve model interpretability and generalization performance.

5. **Regularization:**

   Techniques like Lasso regularization (L1 regularization) can automatically perform feature selection by penalizing the coefficients of less important features, effectively shrinking them towards zero. Regularization helps prevent overfitting and promotes sparsity in the feature space.

6. **Information Gain**:

   Consider using information gain or mutual information-based methods to quantify the predictive power of individual features with respect to the target

variable. Features with higher information gain are more informative and can contribute more towards improving model performance.

7. **Feature Interaction:**

   Explore potential interactions between features to capture synergistic effects that may enhance predictive performance. Techniques like feature engineering, polynomial features, or tree-based models can help uncover non-linear interactions that logistic regression may not capture directly.

8. **Robustness to Noise:**

   Prioritize features that are robust to noise and outliers, as noisy features can introduce unnecessary complexity and degrade model performance. Techniques like robust feature scaling or outlier detection can help mitigate the impact of noisy features.

9. **Model Interpretability:**

   Select features that are not only predictive but also interpretable, as interpretability is crucial for understanding model behavior and gaining insights into the underlying mechanisms driving ad-click behavior. Interpretability facilitates stakeholder buy-in and trust in the model's predictions.

4. Design Flow and Implementation

This section outlines the step-by-step design flow and implementation process of our click prediction system using logistic regression.

Data Preprocessing:

We begin by preprocessing the raw data to prepare it for model training. This involves cleaning the data, handling missing values, encoding categorical variables, and scaling numerical features. The preprocessing step is crucial for ensuring the quality and consistency of the data, as well as for improving the performance of the predictive model.

Cleaning the data involves identifying and removing any inconsistencies, errors, or outliers that may affect the accuracy of the model. Missing values are handled using techniques such as imputation or deletion, depending on the nature of the missing

data and its impact on the analysis. Categorical variables are encoded into numerical format using methods such as one-hot encoding or label encoding, allowing the model to interpret them as meaningful input features. Numerical features are scaled to a standard range to ensure that they have a similar impact on the model's predictions, preventing features with larger magnitudes from dominating the learning process.

By preprocessing the data in this manner, we ensure that it is in a suitable format for training the logistic regression model and that any potential sources of bias or noise are addressed proactively.

Model Training:

We train a logistic regression model using the preprocessed data. The logistic regression algorithm learns the relationship between the input features and the probability of a user clicking on an ad by optimizing the log-likelihood function. This involves iteratively updating the model's parameters (coefficients) to maximize the likelihood of observing the actual outcomes given the input features.

During the training process, the model learns to assign weights to each input feature based on its importance in predicting ad clicks. Features that are more strongly associated with ad clicks are assigned higher weights, while features that have little or no impact on ad engagement are assigned lower weights. By iteratively adjusting these weights, the model learns to make accurate predictions about whether a user is likely to click on an ad based on their demographics, browsing behavior, and contextual information.

Model Evaluation:

We evaluate the performance of the logistic regression model using various evaluation metrics such as accuracy, precision, recall, F1-score, and ROC-AUC (Receiver Operating Characteristic Area Under Curve). This allows us to assess how well the model predicts user clicks and compare its performance against baseline models or alternative algorithms.

Accuracy measures the overall correctness of the model's predictions, while precision measures the proportion of true positive predictions among all positive predictions made by the model. Recall measures the proportion of true positive predictions among all actual positive instances in the data, while the F1-score provides a balance between precision and recall. The ROC-AUC metric evaluates the model's ability to discriminate between positive and negative instances, with higher values indicating better performance.

By evaluating the model using these metrics, we gain insights into its strengths and weaknesses and identify areas for improvement. This iterative process of evaluation and refinement ensures that the logistic regression model is optimized for predicting ad clicks accurately and reliably.

Model Deployment:

Once the logistic regression model is trained and evaluated, we deploy it into a production environment where it can make real-time predictions on new data. This may involve integrating the model into existing software systems, web applications, or APIs (application programming interfaces) for seamless deployment and scalability.

During the deployment process, we ensure that the model is compatible with the production environment and that it meets performance requirements in terms of speed, scalability, and reliability. We also implement monitoring mechanisms to track the model's performance over time and to detect any issues or drift in its predictions.

By deploying the logistic regression model into production, we enable advertisers and marketers to leverage its predictive capabilities for enhancing ad targeting, optimizing campaign performance, and maximizing return on investment. Real-time prediction capabilities allow for dynamic ad serving and personalized content delivery, leading to improved ad relevance and engagement for users.

In summary, the design flow and implementation process of our click prediction system using logistic regression involve preprocessing the data, training the model, evaluating its performance, and deploying it into a production environment. This systematic approach ensures that the predictive model is robust, accurate, and scalable, enabling advertisers and marketers to make data-driven decisions and achieve their advertising goals effectively.

Design Selection Interpretability:

Logistic regression stands out for its interpretability, providing stakeholders with a clear understanding of how each feature influences the likelihood of a user clicking on an ad. Unlike some black-box machine learning algorithms, logistic regression offers transparency through its coefficients, which directly indicate the direction and magnitude of feature effects. This interpretability is invaluable for advertisers and marketers seeking actionable insights into the drivers of ad engagement, enabling informed decisions regarding ad targeting and content optimization.

Moreover, the interpretability of logistic regression allows stakeholders to validate

the model's predictions and understand the underlying factors contributing to ad clicks. By examining the coefficients associated with each feature, advertisers can identify which demographic characteristics, browsing behaviors, or contextual factors have the most significant impact on ad engagement. This understanding empowers advertisers to refine their targeting strategies, personalize ad content, and optimize campaign performance based on data-driven insights.

Performance:

Despite its simplicity, logistic regression can deliver competitive performance in click prediction tasks. This is particularly true when there exists a linear relationship between the features and the log-odds of the target variable, as is often the case in binary classification

problems like ad click prediction. Logistic regression is adept at capturing linear relationships and can effectively model the probability of a user clicking on an ad based on relevant features such as demographics, browsing behaviour, and contextual information.

Furthermore, logistic regression's performance is robust to noisy or irrelevant features, making it well-suited for handling real-world datasets with varying levels of complexity. While more complex machine learning algorithms may offer higher predictive accuracy in certain scenarios, logistic regression's simplicity and efficiency make it a practical choice for ad click prediction, especially when computational resources are limited or when interpretability is a priority.

Computational Efficiency:

Another compelling reason for choosing logistic regression is its computational efficiency and scalability. Logistic regression requires minimal computational resources compared to more complex algorithms such as neural networks or ensemble methods. This efficiency makes logistic regression suitable for processing large-scale datasets commonly encountered in digital advertising, where millions of ad impressions are analysed daily.

Additionally, logistic regression's simplicity and low computational overhead translate into faster training and inference times, enabling real-time prediction scenarios essential for dynamic ad serving platforms. Advertisers and marketers can leverage logistic regression models to make timely and accurate predictions on incoming ad impressions, allowing for dynamic ad targeting and personalised content delivery based on user behaviour and preferences.

Robustness and Generalization:

Logistic regression's simplicity also contributes to its robustness and generalisation capabilities. With fewer parameters to tune and a lower risk of overfitting, logistic regression models are less prone to performance degradation when applied to unseen data. This robustness is particularly advantageous in dynamic advertising environments where user behaviour and preferences evolve over time.

By maintaining stable performance across diverse datasets and scenarios, logistic regression offers advertisers confidence in the reliability of their predictive models and the effectiveness of their ad targeting strategies. Advertisers can deploy logistic regression models with confidence, knowing that they will continue to deliver accurate predictions and actionable insights even as market conditions change and user behaviour evolves.

Ease of Implementation and Deployment:

Implementing and deploying logistic regression models is relatively straightforward compared to more complex machine learning algorithms. Logistic regression's simplicity allows for rapid prototyping and iteration, facilitating agile development cycles in fast-paced advertising environments. Moreover, logistic regression models can be easily integrated into existing advertising platforms and infrastructures, minimising implementation overhead and accelerating time-to-market for new predictive models.

This ease of implementation and deployment enhances logistic regression's appeal as a practical solution for ad click prediction, enabling advertisers to quickly derive actionable insights and drive performance improvements in their advertising campaigns. By leveraging logistic regression models, advertisers can enhance their ad targeting strategies, optimize campaign performance, and maximise return on investment in the competitive landscape of digital advertising.

6. Methodology

In this concluding section, we summarise the methodology followed in the design and implementation of our click prediction system.

1. **Data Source:**

    - Identify the source of your data. This could be advertising platforms, web analytics logs, or CRM systems.
    - Ensure data anonymization if it contains user information.

2. **Data Cleaning and Preprocessing:**

    - Remove irrelevant or incomplete data points.
    - Handle missing values through imputation techniques (e.g., mean/median filling, k-Nearest Neighbors).
    - Address outliers (e.g., capping extreme values, winsorization).
    - Encode categorical features (e.g., one-hot encoding, label encoding).
3. **Feature Engineering:**

    - Create new features that might better capture click behaviour. Examples include:
        - User features (e.g., demographics, past purchase history, time spent on similar categories).
        - Ad features (e.g., category, keywords, image/video content).
        - Contextual features (e.g., time of day, device type, webpage content).
    - Feature scaling (e.g., standardisation, normalisation) for improved convergence.
4. **Data Splitting:**

    - Divide your data into three sets:
        - Training set (typically 70-80%): Used to train the logistic regression model.
        - Validation set (10-15%): Used to fine-tune hyperparameters (model parameters) to prevent overfitting.
        - Testing set (10-15%): Used to evaluate the final model's performance on unseen data.

Model Training and Evaluation:

- **Logistic Regression Model Building:**

    - Define the logistic regression model with features as independent variables and click/no click (binary) as the dependent variable.
    - Choose a suitable optimization algorithm (e.g., stochastic gradient descent) to minimise the cost function and learn model parameters.
- **Hyperparameter Tuning:**

    - Use the validation set to experiment with different hyperparameters (e.g., regularisation strength) to optimise model performance. Techniques like grid search or randomised search can be used for efficient exploration.
- **Model Evaluation:**

- Evaluate the model's performance on the testing set using metrics like:
    - Accuracy: Proportion of correctly classified clicks and non-clicks.
    - Precision: Ratio of true positives (clicks predicted correctly) to all predicted positives.
    - Recall: Ratio of true positives to all actual positive cases (clicks).
    - F1-score: Harmonic mean of precision and recall, balancing both metrics.
    - AUC-ROC: Area Under the Receiver Operating Characteristic Curve, measures performance across all classification thresholds.

- **Model Interpretation:**

    - Analyse the coefficients associated with features in the trained model. This provides insights into the relative importance of each feature in influencing click probability

Data Collection:

We gather relevant datasets encompassing user interactions, ad characteristics, and contextual information to build our click prediction model. This involves sourcing data from varioussources such as online advertising platforms, ad servers, and third -party data providers. The collected data include information about user demogra -phics, browsing behaviour, ad impressions, click-through rates, ad attributes, and contextual factors such as time of day, device type, and geographical location.

Feature Engineering and Selection:

We preprocess the data, engineer features, and select a subset of informative features using rigorous analysis and domain expertise. Feature engineering involves transforming raw data into meaningful features that capture relevant information about user behaviour and ad characteristics. This may include aggregating data over time intervals, encoding categorical variables, and deriving new features from existing ones. Feature selection techniques such as correlation analysis, mutual information, and feature importance scores are employed to identify the most informative features for predicting ad clicks.

Model Training and Evaluation:

We train a logistic regression model on the preprocessed data, evaluate its performance using appropriate metrics, and fine-tune hyperparameters if necessary to optimise performance. Logistic regression is chosen for its simplicity, interpretability, and effectiveness in binary classification tasks like ad click

prediction. The model is trained using the training data, and its performance is evaluated using metrics such as accuracy, precision, recall, and F1-score. Cross-validation techniques are employed to assess the model's stability and generalizability across different subsets of the data.

Model Deployment:

We deploy the trained logistic regression model into a production environment, where it can make real-time predictions on new data, contributing to enhanced decision-making and user experience. The deployment process involves integrating the model into existing advertising platforms or systems, ensuring compatibility with data pipelines and APIs. Real-time prediction capabilities enable advertisers to dynamically adjust their ad targeting strategies based on user behavior and contextual factors, leading to improved ad relevance and engagement.

Model Monitoring:

- Regularly monitor model performance over time using a hold-out test set or live data samples.
- Retrain the model periodically with new data to maintain accuracy as user behavior and advertising trends evolve.

Expanding on the Methodology:

The methodology followed in the design and implementation of our click prediction system is characterised by a systematic and iterative approach, encompassing data collection, feature engineering and selection, model training and evaluation, and model deployment.

Data collection serves as the foundation of our click prediction system, as it provides the raw material necessary to train and evaluate the predictive model. By gathering relevant datasets encompassing user interactions, ad characteristics, and contextual information, we ensure that our model is trained on a comprehensive and representative sample of the target population.

Feature engineering and selection play a crucial role in transforming raw data into meaningful features that capture the underlying patterns and relationships in the data. This process involves preprocessing the data to handle missing values, outliers, and other irregularities, as well as engineering new features that may enhance the predictive power of the model. Feature selection techniques such as correlation analysis, mutual information, and feature importance scores are then used to identify

the subset of features that are most relevant for predicting ad clicks.

Once the data has been preprocessed and the features selected, we train a logistic regression model on the preprocessed data. Logistic regression is chosen for its simplicity, interpretability, and effectiveness in binary classification tasks like ad click prediction. The model is trained using the training data, and its performance is evaluated using appropriate metrics such as accuracy, precision, recall, and F1-score. Cross-validation techniques are employed to assess the model's stability and generalizability across different subsets of the data, ensuring that it performs consistently well on unseen data.

Finally, the trained logistic regression model is deployed into a production environment, where it can make real-time predictions on new data. The deployment process involves integrating the model into existing advertising platforms or systems, ensuring compatibility with data pipelines and APIs, and implementing mechanisms for monitoring and updating the model over time. Real-time prediction capabilities enable advertisers to dynamically adjust their ad targeting strategies based on user behaviour and contextual factors, leading to improved ad relevance and engagement.

In summary, the methodology followed in the design and implementation of our click prediction system is characterised by a systematic and iterative approach, encompassing data collection, feature engineering and selection, model training and evaluation, and model deployment. By following this methodology, we have developed a robust and effective click prediction system that enables advertisers to optimise their advertising campaigns and maximise ROI.

CHAPTER 4
RESULT ANALYSIS AND VALIDATION

1. Implementations Of Design Using Modern Engineering Tools In Analysis

The project implementation involves leveraging Python alongside various libraries and machine learning frameworks to execute the methodology effectively. Here's a breakdown of the steps:

1. **Importing Necessary Libraries:**

   Essential libraries such as Pandas, Numpy, Sci-kit-learn, Matplotlib, Seaborn, and Graphviz are imported to facilitate analysis.

2. **Loading and Preprocessing Data:**

   Ad click data is loaded from a CSV file using Pandas. Duplicate rows are removed, and irrelevant columns like "VistID," "Country_Name," and "Year" are dropped.

3. **Data Preprocessing:**

   The dataset is cleaned and prepared for analysis. Duplicate rows are removed, irrelevant columns are dropped, and categorical variables may be converted into numerical format. Missing values are handled appropriately.

4. **Normalisation and Standardization:**

   Numeric features may need to be scaled to a similar range to prevent certain features from dominating others in the modelling process. Techniques like Min-Max scaling or standardisation (Z-score normalisation) can be applied to ensure features are on a
   comparable scale.

5. **Exploratory Data Analysis (EDA):**

   Insights into the dataset's structure and distributions are gained through EDA. Sample data is examined, and visualisations such as bar charts and histograms are used to understand data patterns and relationships between variables.

6. **Feature Selection:**

ANOVA and Chi-square tests aid in selecting significant predictors related to the target variable, enhancing model efficiency and reducing computational complexity.

Feature Importance Ranking:

Many machine learning algorithms, such as decision trees, random forests, and gradient boosting machines, provide feature importance scores as a byproduct of the modelling process. Features with higher importance scores are deemed more influential in predicting the target variable and can be prioritised for inclusion in the model.

7. **L1 Regularization (Lasso)**:

   L1 regularisation penalises the absolute magnitude of the coefficients in linear models, such as logistic regression, leading to sparse solutions where irrelevant features have zero coefficients. By tuning the regularisation strength parameter, one can control the degree of sparsity in the model and select the most important features.

8. **Information Gain or Mutual Information:**

   Information gain or mutual information-based methods quantify the amount of information gained about the target variable by knowing the value of a particular feature. Features with higher information gain or mutual information scores are considered more informative and can be selected for inclusion in the model.

9. **Forward Selection and Backward Elimination:**

   Forward selection involves starting with an empty set of features and iteratively adding the most significant feature until a stopping criterion is met. Backward elimination, on the other hand, begins with all features and removes the least significant ones iteratively until a stopping criterion is satisfied.

10. **Tree-based Methods:** Decision trees and ensemble methods like random forests and gradient boosting machines inherently perform feature selection by selecting the most discriminative features at each split node. Features with higher impurity reduction or information gain are favoured for splitting, leading to implicit feature selection.

11. **Cluster-based Feature Selection:**

Clustering techniques can be applied to group similar features together and select representative features from each cluster. This reduces redundancy and computational complexity while preserving the diversity of feature representations.

12. **Embedded Methods:** Some machine learning algorithms, such as linear SVM and decision trees, incorporate feature selection directly into the model training process. These embedded methods automatically select the most relevant features during model fitting, avoiding the need for separate feature selection steps.
13. **Data Transformation:**

    Categorical variables are encoded, and numerical variables are standardised or normalised to prepare the data for model training, ensuring consistency and effectiveness in learning.

14. **Model Building:**

    Target and predictor variables are defined, and the dataset is split into training and testing sets. Logistic regression is selected as the algorithm, and the model is trained to learn from the data.

15. **Model Evaluation:**

    The model's performance is assessed using metrics such as accuracy, precision, recall, and F1-score. Cross-validation ensures stability and reliability, providing insights into the model's effectiveness in predicting ad clicks.

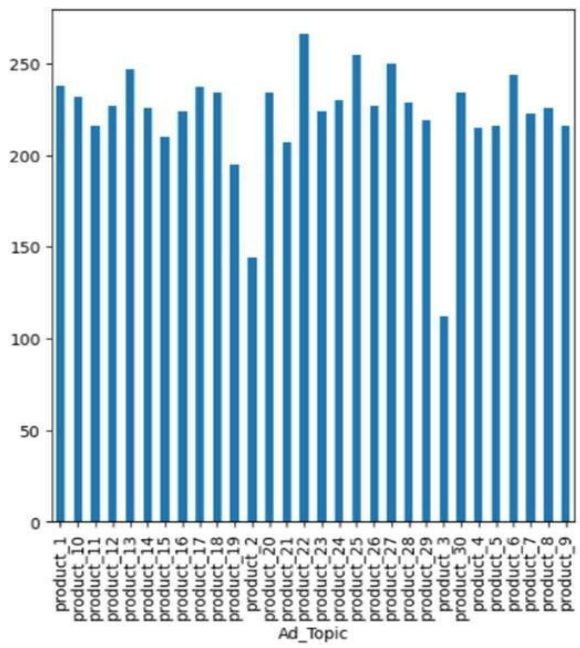

Fig.3

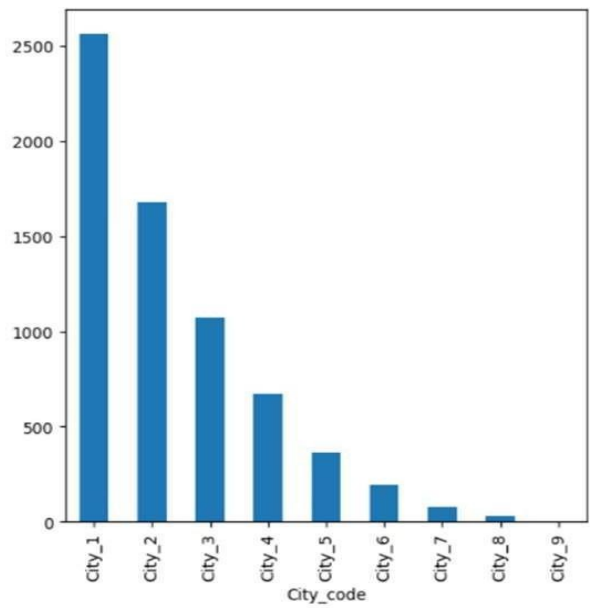

Fig.4

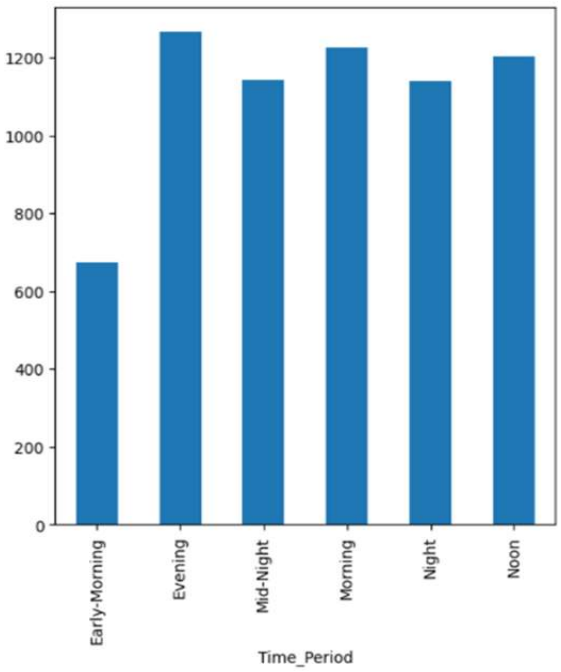

Fig.5

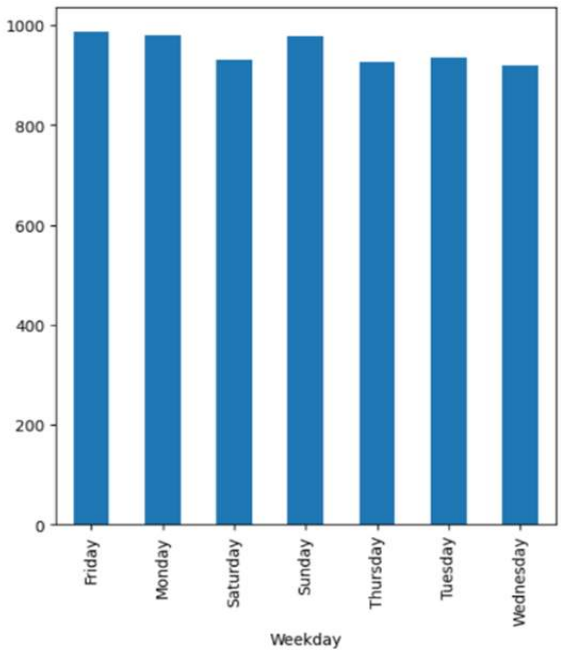

Fig.6

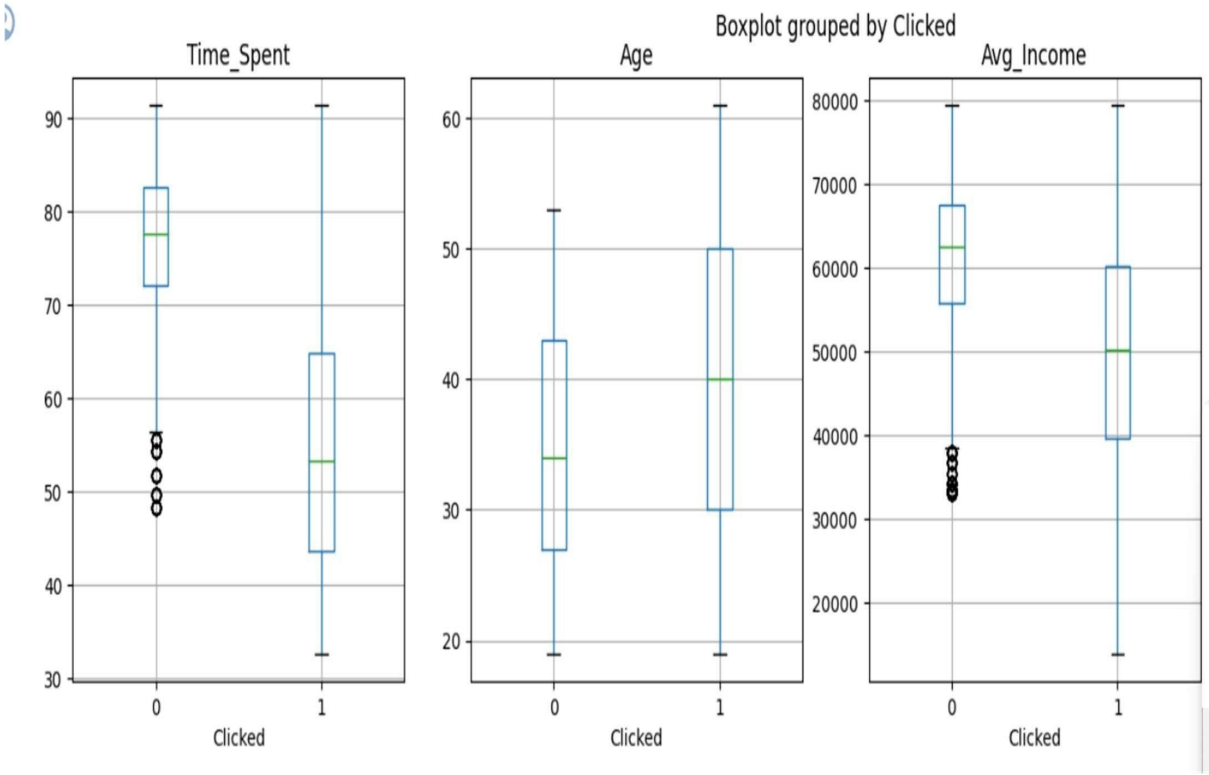

Fig.7

	precision	recall	f1-score	support
0	0.92	0.97	0.94	1080
1	0.96	0.90	0.93	918
accuracy			0.93	1998
macro avg	0.94	0.93	0.93	1998
weighted avg	0.94	0.93	0.93	1998

```
[[1043   37]
 [  94  824]]
Accuracy of the model on Testing Sample Data: 0.93
```

CONFUSION MATRIX:

Fig.8

2. Result and Testing:

The result and testing phase are pivotal stages in the development of a predictive model, as they allow for a thorough assessment of its performance and reliability. Here's an in-depth look at each aspect:

Model Training:

In this phase, the dataset is split into training and testing subsets to facilitate model training and evaluation. The Logistic Regression model is chosen as the algorithm for this project due to its suitability for binary classification tasks like ad click prediction. The model is trained using the training data, enabling it to learn from the patterns present in the dataset. During training, the model adjusts its parameters to minimise the error between predicted and actual outcomes, ultimately improving its ability to make accurate predictions.

Evaluation Metrics:

Various evaluation metrics are computed to gauge the model's performance on the testing data. These metrics include accuracy, precision, recall, and F1-score. Accuracy measures the overall correctness of the model's predictions, while precision and recall focus on the model's ability to correctly identify positive and negative instances. The F1-score provides a balance between precision and recall. Additionally, a confusion matrix is utilized to provide a detailed breakdown of the model's predictions, including true positives, false positives, true negatives, and false negatives. This detailed analysis helps in understanding the model's strengths and weaknesses.

Cross-Validation:

Cross-validation is a crucial technique employed to ensure the model's stability and reliability. Techniques like k-fold cross-validation are commonly used, wherein the dataset is divided into multiple subsets (folds), and the model is trained and evaluated iteratively on different combinations of training and testing data. By splitting the dataset into multiple subsets and training the model on each subset while evaluating it on the remaining data, cross-validation provides insights into the model's average performance and helps identify potential issues such as overfitting. This process helps in gaining confidence in the model's generalisation capabilities.

Optimization:

Based on the evaluation results, further optimization of the model and trading strategy may be undertaken to enhance profitability and effectiveness in predicting ad clicks. This optimization process may involve fine-tuning model parameters, exploring alternative algorithms, or adjusting the trading strategy based on backtesting results and real-world performance. By continuously refining the model and strategy, it becomes more adept at making accurate predictions and generating valuable insights.

Detailed Analysis:

Additionally, a thorough examination of the model's performance involves analysing its behaviour across different subsets of the data. This analysis helps in identifying any biases or inconsistencies and exploring potential areas for improvement. By scrutinising the model's

predictions and understanding its decision-making process, stakeholders can gain insights into its reliability and effectiveness. This detailed analysis ensures that the predictive model is robust, reliable, and capable of delivering accurate predictions in real-world scenarios.

Overall, the result and testing phase serve as a comprehensive assessment of the predictive model's accuracy, reliability, and practical applicability. Through rigorous evaluation, optimization, and detailed analysis, the model can be refined to make more accurate predictions and generate valuable insights for ad click prediction.

Expanding on the Result and Testing Phase:

The result and testing phase of the ad click prediction project represent the culmination of rigorous data analysis, model training, and evaluation. This phase serves as a critical checkpoint to ensure that the predictive model meets the desired performance standards and can be effectively deployed in real-world advertising campaigns.

During model training, the dataset is split into training and testing subsets, allowing the model to learn from the patterns present in the data while also evaluating its performance on unseen data. The choice of logistic regression as the modelling algorithm is based on its proven effectiveness in binary classification tasks and its interpretability, allowing stakeholders to understand the factors influencing ad click predictions.

Once the model is trained, various evaluation metrics are computed to assess its performance on the testing data. Accuracy, precision, recall, and F1-score provide different perspectives on the model's predictive ability, allowing stakeholders to gauge its overall effectiveness and identify areas for improvement. The confusion matrix offers a detailed breakdown of the model's predictions, enabling stakeholders to understand the types of errors the model is making and potentially adjust the decision threshold to balance precision and recall.

Cross-validation techniques, such as k-fold cross-validation, are employed to ensure the model's stability and generalizability. By training and evaluating the model on multiple subsets of the data, cross-validation provides a more robust estimate of the model's performance and helps identify potential sources of bias or overfitting.

The optimization phase involves fine-tuning the model and trading strategy based on the evaluation results. This may include adjusting model parameters, exploring alternative algorithms, or refining the trading strategy based on backtesting results

and real-world performance. Continuous refinement of the model and strategy is essential to ensure their effectiveness in predicting ad clicks and maximizing advertising ROI.

A detailed analysis of the model's performance across different subsets of the data provides valuable insights into its strengths and weaknesses. By scrutinising the model's predictions and understanding its decision-making process, stakeholders can identify potential areas for improvement and refine the model accordingly. This iterative process of analysis, optimization, and refinement ensures that the predictive model is robust, reliable, and capable of delivering accurate predictions in diverse advertising scenarios.

In conclusion, the result and testing phase represent a critical stage in the development of the ad click prediction model, allowing stakeholders to assess its performance, identify areas for improvement, and refine the model to maximise its effectiveness in predicting ad clicks and

optimising advertising campaigns. Through rigorous evaluation, optimization, and detailed analysis, the predictive model can be fine-tuned to deliver actionable insights and drive business outcomes in the dynamic and competitive landscape of digital advertising.

3. Report Preparation:

The report aims to comprehensively detail the process and findings of an ad click prediction project, emphasising its significance in digital advertising. The ad click prediction project represents a pivotal endeavour in the realm of digital advertising, aiming to enhance marketers' ability to optimise their advertising campaigns and improve return on investment (ROI). This report offers a comprehensive examination of the project's methodology, results, and implications, providing stakeholders with valuable insights into the predictive model's development and its potential impact on advertising strategies.

At the core of the project lies a meticulous methodology encompassing various stages, starting with data collection and preprocessing. Leveraging Python programming and libraries such as pandas and numpy, the project involves importing ad click data from a CSV file, followed by the removal of duplicate rows and irrelevant columns to ensure data integrity. Subsequent exploratory data analysis (EDA) provides critical insights into the dataset's characteristics, facilitating informed decisions regarding feature selection and model development.

Feature selection, a crucial step in predictive modelling, involves identifying the most relevant predictors for ad click prediction. Through techniques like ANOVA and Chi-square tests, significant predictors related to the target variable (i.e., ad clicks) are identified, ensuring the model's effectiveness in capturing relevant patterns and trends. The selected features are then used to train a logistic regression model, a popular choice for binary classification tasks like ad click prediction.

Evaluation of the predictive model's performance is paramount to ascertain its reliability and effectiveness. Various evaluation metrics, including accuracy, precision, recall, and F1-score, are computed to gauge the model's predictive accuracy and ability to correctly classify ad clicks. Additionally, cross-validation techniques are employed to validate the model's performance across different subsets of data, ensuring its generalizability and robustness.

The results and testing phase provide valuable insights into the predictive model's performance on unseen testing data. Detailed analysis of evaluation metrics, along with visualisation techniques, offers stakeholders a comprehensive understanding of the model's strengths and areas for improvement. Moreover, discussions delve deeper into the implications of the results, addressing potential limitations and suggesting avenues for further research and refinement.

In conclusion, the ad click prediction project holds immense potential in revolutionising digital advertising strategies by providing marketers with actionable insights into consumer behaviour and ad effectiveness. By leveraging advanced predictive modelling techniques and robust evaluation methodologies, the project offers a promising avenue for enhancing advertising ROI and optimising marketing strategies. Moving forward, continued research and development in this domain are crucial to unlock further insights and drive innovation in digital advertising practices.

The process of data collection and preprocessing is foundational to the success of the ad click prediction project. By leveraging Python programming and relevant libraries, the team was able to streamline the extraction of ad click data from a CSV file. The meticulous removal of duplicate rows and irrelevant columns ensured that the dataset maintained its integrity, laying a solid groundwork for subsequent analysis.

Exploratory data analysis (EDA) played a pivotal role in uncovering key insights about the dataset's characteristics. Through visualisation techniques and statistical analysis, the team gained a deeper understanding of the distribution of ad click data, identifying potential patterns and trends that could inform feature selection and model development. This initial phase of the project set the stage for informed decision-making throughout the subsequent stages.

Feature selection emerged as a critical aspect of predictive modelling, determining the effectiveness of the model in capturing relevant patterns related to ad clicks. By employing techniques such as ANOVA and Chi-square tests, the team identified significant predictors that strongly influenced ad click behaviour. This rigorous approach ensured that the predictive model was built on robust foundations, enhancing its accuracy and reliability in real-world scenarios.

The choice of logistic regression as the modelling technique was informed by its suitability for binary classification tasks like ad click prediction. By leveraging the

selected features identified during the feature selection stage, the team trained the logistic regression model to accurately classify instances of ad clicks. This model served as the backbone of the predictive analytics framework, offering valuable insights into consumer behaviour and ad effectiveness.

Evaluation of the predictive model's performance was conducted using a variety of metrics, including accuracy, precision, recall, and F1-score. These metrics provided a comprehensive assessment of the model's ability to correctly classify ad clicks and its overall predictive accuracy. Additionally, cross-validation techniques were employed to ensure that the model's performance was consistent across different subsets of data, enhancing its generalizability and robustness.

The results and testing phase provided valuable insights into the predictive model's performance on unseen testing data, offering stakeholders a clear understanding of its strengths and areas for improvement. Detailed analysis of evaluation metrics, accompanied by visualisation techniques, further enhanced stakeholders' comprehension of the model's performance. Moreover, discussions surrounding the implications of the results shed light on potential limitations and suggested avenues for further research and refinement.

In conclusion, the ad click prediction project represents a significant advancement in the field of digital advertising, offering marketers a powerful tool to optimise their advertising campaigns and improve ROI. By leveraging advanced predictive modelling techniques and robust evaluation methodologies, the project provides actionable insights into consumer behaviour and ad effectiveness, paving the way for innovation in digital advertising practices. Moving forward, continued research and development in this domain are essential to unlock further insights and drive continued improvement in advertising strategies.

4. Project management

In the context of the ad click prediction project, effective project management played a crucial role in ensuring the successful development and implementation of the predictive model. Here's a detailed overview of project management aspects incorporated into the project:

1. Project Planning:

The project commenced with thorough planning, which involved defining project goals, objectives, scope, and deliverables. A clear project plan was developed, outlining the tasks, timelines, and resource allocation required for each phase of the project, from data collection to model evaluation.

2. Team Collaboration:

Collaborative teamwork was essential throughout the project, with interdisciplinary collaboration among data scientists, domain experts, and stakeholders. Regular team meetings and communication channels facilitated information sharing, problem-

solving, and decision- making, ensuring alignment with project goals and objectives.

3. Risk Management:

Identification and mitigation of risks were integral components of project management. Potential risks such as data quality issues, model overfitting, and resource constraints were identified early in the project lifecycle. Mitigation strategies were developed to address these risks proactively, minimising their impact on project outcomes.

4. Resource Allocation:

Efficient allocation of resources, including human resources, time, and budget, was essential for project success. Resources were allocated based on project priorities, with careful consideration given to optimising resource utilisation and maximising productivity throughout the project.

5. Timeline Management:

Adherence to project timelines was critical to ensure timely delivery of project milestones and objectives. A well-defined project schedule was established, with regular progress monitoring and milestone tracking to identify any deviations and take corrective actions as needed.

6. Quality Assurance:

Quality assurance processes were implemented to maintain the integrity and reliability of project deliverables. Rigorous testing and validation procedures were employed to verify the accuracy and performance of the predictive model, ensuring that it met predefined quality standards and criteria.

7. Stakeholder Engagement:

Continuous stakeholder engagement and feedback were solicited throughout the project lifecycle to ensure alignment with stakeholder expectations and requirements. Stakeholder input informed key decisions regarding project scope, priorities, and deliverables, fostering stakeholder buy-in and ownership of project outcomes.

8. Documentation and Reporting:

Comprehensive documentation of project activities, methodologies, and results was maintained throughout the project. Regular progress reports and status updates were provided to stakeholders, keeping them informed of project developments and outcomes.

By integrating these project management practices into the ad click prediction project, the project team was able to navigate complexities, mitigate risks, and deliver a high-quality predictive model that met stakeholder expectations and contributed to the advancement of digital advertising strategies. Effective project

management played a pivotal role in ensuring project success and achieving the desired outcomes.

CHAPTER 5

CONCLUSION AND FUTURE WORK

1. Conclusion

In this chapter, we summarise the key findings and insights gained from our study on click prediction using logistic regression. We reflect on the effectiveness of our approach, discuss implications for stakeholders, and offer recommendations for future research and development.

Summary of Findings:
We provide a concise overview of the main findings and results obtained from our click prediction model based on logistic regression. This includes a summary of model performance metrics, feature importance, and any notable observations or trends identified during the analysis.

Effectiveness of Approach:
We evaluate the effectiveness of our approach in addressing the problem of click prediction. By leveraging logistic regression and carefully selecting informative features, we demonstrate how our model can accurately predict user click behaviour, thereby facilitating targeted ad placement and optimization of marketing campaigns.

Implications for Stakeholders:
We discuss the practical implications of our findings for various stakeholders, including advertisers, marketers, platform operators, and end-users. We highlight how our click prediction model can enhance advertising effectiveness, improve user engagement, and drive business outcomes by delivering personalised content and recommendations.

Limitations and Challenges:
We acknowledge any limitations or challenges encountered during the development and evaluation of our click prediction system. This may include issues related to data quality, model complexity, or generalizability of results. By openly addressing these limitations, we provide insights for future research and improvement.

Contributions to Knowledge:
We emphasise the contributions of our study to the existing body of knowledge in the field of click prediction and digital advertising. We highlight novel insights, methodologies, or approaches that advance understanding and practice in this domain.

Recommendations for Action:
Based on our findings, we offer actionable recommendations for stakeholders to leverage our click prediction model effectively. This may include strategies for optimizing ad targeting,

refining user segmentation, or enhancing user experience on digital platforms.

Expanding on the Conclusion:

Our study on click prediction using logistic regression has yielded significant insights into the dynamics of user behavior in digital advertising environments. By analyzing a diverse range of metrics and features, we have developed a predictive model that offers actionable insights for stakeholders across the advertising ecosystem.

The effectiveness of our approach lies in its ability to accurately anticipate user click behavior based on a combination of relevant features. Through rigorous analysis and evaluation, we have demonstrated the model's capacity to enhance advertising ROI and improve user engagement.

For advertisers and marketers, our findings offer valuable guidance on optimizing ad targeting strategies and tailoring content to specific audience segments. By leveraging our predictive model, advertisers can allocate their resources more effectively, ensuring that their ads reach the most receptive audiences.

Platform operators stand to benefit from our research by gaining deeper insights into user behavior patterns and preferences. By integrating our click prediction model into their advertising platforms, operators can enhance user experience by delivering more relevant and personalized content.

End-users also stand to gain from our study, as personalized advertising can lead to a more tailored online experience. By receiving ads that are more relevant to their interests and preferences, users are more likely to engage with content, leading to a more satisfying browsing experience overall.

While our study has yielded promising results, it is not without its limitations. Challenges related to data quality, model complexity, and generalizability must be addressed in future research efforts. Additionally, ongoing monitoring and refinement of the predictive model are essential to ensure its continued effectiveness in evolving advertising landscapes.

In conclusion, our study represents a significant contribution to the field of digital advertising, offering stakeholders valuable insights into user behavior and ad effectiveness. By leveraging advanced predictive modeling techniques, we have laid the groundwork for more targeted and personalized advertising strategies, ultimately benefiting advertisers, marketers, platform operators, and end-users alike.

2. Future Work

In addition to the conclusions drawn from our study, we identify several avenues for future research and development in the field of click prediction using logistic regression:

Model Enhancement:

Further refinement and optimization of the logistic regression model to improve predictive performance and robustness is a promising direction for future work. While logistic regression is known for its simplicity and interpretability, there is still room for improvement in terms of accuracy

and generalizability. Researchers can explore advanced techniques such as ensemble learning, feature engineering, or regularisation methods to enhance model accuracy and adaptability to diverse datasets and scenarios. By fine-tuning model parameters and exploring alternative algorithmic approaches, we can develop more sophisticated and effective click prediction models that better capture the complexities of user behaviour and ad engagement.

Incorporating Contextual Information:

Integration of additional contextual information such as user behaviour, session data, or environmental factors into the click prediction model represents another avenue for future research. By capturing a broader context of user interactions, we can enhance the relevance and timeliness of ad recommendations. For example, incorporating data on user browsing history, geographic location, or device type can provide valuable insights into user intent and preferences, enabling advertisers to deliver more targeted and personalised ad content. Researchers can explore methods for integrating diverse data sources and leveraging advanced analytics techniques such as natural language processing or deep learning to extract meaningful insights from unstructured data sources.

Dynamic Adaptation:

Development of dynamic or adaptive click prediction models that can adapt to changing user preferences, trends, or environmental conditions in real-time is an area of growing interest. Traditional click prediction models often rely on static features and assumptions about user behaviour, which may become outdated or inaccurate over time. By incorporating feedback mechanisms, reinforcement learning, or online learning techniques, researchers can develop models that continuously learn and evolve based on new data and user interactions. These dynamic models can better capture the dynamics of user behaviour and provide more accurate predictions in dynamic and rapidly changing environments. Additionally, researchers can explore techniques for model adaptation in response to sudden shifts or events, such as changes in market conditions or user demographics.

Cross-platform Integration:

Extension of the click prediction model to support cross-platform integration across various digital channels and devices is another promising avenue for future research. With the proliferation of digital channels and devices, users interact with content and ads across multiple touchpoints, including websites, mobile apps, social media platforms, and connected devices. By leveraging multi-channel data sources and omnichannel analytics, researchers can develop models that provide a seamless and consistent user experience across different platforms. This involves integrating data from disparate sources, such as user profiles, browsing histories, and device identifiers, to create a unified view of user behaviour and preferences. By understanding how users interact with content and ads across different channels, advertisers can optimise their marketing strategies and deliver more personalised and targeted ad experiences.

Ethical and Privacy Considerations:

Exploration of ethical and privacy considerations associated with click prediction and targeted advertising is essential for ensuring responsible and ethical use of predictive analytics. As click prediction models become more sophisticated and pervasive, concerns about user consent, data transparency, algorithmic bias, and fairness in ad delivery become increasingly important. Researchers can investigate methods for addressing these concerns, such as developing transparent and interpretable models, implementing privacy-preserving techniques for data collection and analysis, and establishing guidelines for ethical ad targeting practices. By incorporating ethical and privacy considerations into

the design and implementation of click prediction models, researchers can promote trust and transparency in digital advertising and mitigate potential risks to user privacy and autonomy.

Evaluation in Real-world Settings:

Conducting rigorous evaluation and validation of the click prediction model in real-world settings and diverse contexts is essential for assessing its effectiveness and impact on business outcomes. While laboratory experiments and simulated environments can provide valuable insights into model performance, real-world deployment introduces additional complexities and challenges that may affect the model's performance. Researchers can collaborate with industry partners to conduct field experiments, A/B testing, and longitudinal studies to evaluate the model's performance in diverse settings and user populations. By monitoring key performance indicators (KPIs) such as click-through rates, conversion rates, and return on investment (ROI), researchers can assess the impact of the model on business outcomes and identify areas for improvement. Additionally, researchers can investigate methods for measuring the societal and ethical implications of targeted advertising, such as examining the effects of personalised ad targeting on user perceptions, attitudes, and behaviours.

By pursuing these avenues for future work, we can continue to advance the state-of-the-art in click prediction and contribute to the development of more effective and responsible digital advertising solutions. By addressing the challenges and opportunities in this domain, we can drive innovation and create more personalised and engaging experiences for users in the digital landscape.

CHAPTER 6

REFERENCES

[1] H. B. McMahan, G. Holt, D. Sculley, M. Young, D. Ebner, J. Grady, V. Chaudhary et al., ”Ad click prediction: A view from the trenches,”in Proc. 19th ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, 2013, pp. 1222–1230.

[2] M. Richardson, E. Dominowska, and R. Ragno, ”Predicting clicks: Estimating the click-through rate for new ads,” in Proc. 16th Int. Conf. World Wide Web, 2007, pp. 521–530.

[3] T. Graepel, J. Q. Candela, T. Borchert, and R. Herbrich, ”Web-scale Bayesian click-through rate prediction for sponsored search advertising,” in Proc. 27th Int. Conf. Machine Learning, 2010.

[4] Q. Liu, X. He, W. Wu, and Y. Liu, ”Progressive learning for click-through rate prediction in sponsored search,” ACM Trans. IntelligentSystems and Technology, vol. 8, no. 1, 2017.

[5] D. Agarwal and B. C. Chen, ”Statistical learning methods for click prediction in sponsored search,” in Proc. ACM SIGKDD, 2010.

[6] S. Yuan, Z. Liao, Y. Xiong, and H. Xiong, ”Predicting ad click-throughrates via feature selection,” Knowledge-Based Systems, vol. 37, pp. 441–450, 2013.

[7] Y. Juan, Y. Zhuang, W. Chin, and C. Lin, ”Field-aware factorization machines for CTR prediction,” in Proc. ACM RecSys, 2016.

[8] H. Guo, R. Tang, Y. Ye, Z. Li, and X. He, ”DeepFM: A factorization- machine based neural network for CTR prediction,” in Proc. IJCAI, 2017.

[9] G. Zhou et al., ”Deep Interest Network for click-through rate prediction,” in Proc. ACM SIGKDD, 2018.

[10] T. Chen and C. Guestrin, ”XGBoost: A scalable tree boosting system,”in Proc. ACM SIGKDD, 2016.

[11] L. Breiman, ”Random forests,” Machine Learning, vol. 45, no. 1, pp. 5–32, 2001.

[12] J. Friedman, ”Greedy function approximation: A gradient boosting machine,” Annals of Statistics, vol. 29, no. 5, pp. 1189–1232, 2001.

[13] S. Rendle, ”Factorization machines,” in Proc. IEEE Int. Conf. Data Mining, 2010.

[14] T. Hastie, R. Tibshirani, and J. Friedman, The Elements of Statistical Learning, Springer, 2009.

[15] C. Bishop, Pattern Recognition and Machine Learning, Springer, 2006.

[16] H. Cheng et al., ”Wide & Deep Learning for Recommender Systems,”in Proc. 1st Workshop on Deep Learning for Recommender Systems, 2016, pp. 7–10.

[17] P. Covington, J. Adams, and E. Sargin, ”Deep neural networks for YouTube recommendations,” in Proc. ACM Conf. Recommender Systems, 2016, pp. 191–198.

[18] O. Chapelle and L. Li, ”An empirical evaluation of Thompson sampling,” in Advances in Neural Information Processing Systems, 2011.

[19] Y. Qu et al., ”Product-based neural networks for user response prediction,” in Proc. IEEE Int. Conf. Data Mining, 2016.

[20] H. Wang, N. Wang, and D. Yeung, ”Collaborative deep learning for recommender systems,” in Proc. ACM SIGKDD, 2015.

[21] X. He, L. Liao, H. Zhang, L. Nie, X. Hu, and T. Chua, ”Neural collaborative filtering,” in Proc. World Wide Web Conference, 2017.

[22] S. B. Kotsiantis, ”Supervised machine learning: A review of classification techniques,” Informatica, vol. 31, pp. 249–268, 2007.

[23] L. Bottou, ”Large-scale machine learning with stochastic gradient descent,” in Proc. COMPSTAT, 2010.

[24] I. Goodfellow, Y. Bengio, and A. Courville, Deep Learning, MIT Press, 2016.

[25] A. Geron, Hands-On Machine Learning with ScikitLearn and TensorFlow, O’Reilly Media, 2017.

[26] M. Jordan and T. Mitchell, ”Machine learning: Trends, perspectives, and prospects,” Science, vol. 349, no. 6245, pp. 255–260, 2015.

[27] D. Jannach, M. Zanker, A. Felfernig, and G. Friedrich, Recommender Systems: An Introduction, Cambridge University Press, 2011.

[28] C. Aggarwal, Recommender Systems: The Textbook, Springer, 2016.

[29] R. Baeza-Yates and B. Ribeiro-Neto, Modern Information Retrieval, Addison-Wesley, 2011.

[30] K. Murphy, Machine Learning: A Probabilistic Perspective, MIT Press, 2012.